

# Near-Stream Computing: General and Transparent Near-Cache Acceleration

Zhengrong Wang, Jian Weng, Sihao Liu, Tony Nowatzki

UCLA


April. 2022



# Near-Data Computing to Decentralize Computation

```
while (i < N)
  C[i] = A[i] + B[i]
  i++
```

**Requesting Core**



**Remote L3 Cache**


**Data movement become increasingly the bottleneck.**

- Computation centralized in the core.
- More expensive to fetch data as system scales up.
- Prefetching/streaming can only partially help.

# Near-Data Computing to Decentralize Computation

```
while (i < N)
  C[i] = A[i] + B[i]
  i++
```

Requesting Core



**Data movement become increasingly the bottleneck.**

- Computation centralized in the core.
- More expensive to fetch data as system scales up.
- Prefetching/streaming can only partially help.

**Near-data computing decentralizes the computation.**


- Location: memory, storage, *LLC (Our Focus)*, ...
- Technology: in-situ, inorder cores, FPGA, ...

**Existing NDC works fall short to fully cover:**

- **Transparency.**
- **Synchronization-Efficiency.**
- **Generality.**

**... due to inappropriate *abstraction* choice:**

- Instructions, user-defined functions, threads, ...



# Features of An Ideal NDC Abstraction

## Transparency.


- No manual programming, auto-transformation, sequential semantics.

## Synchronization-Efficiency.

- Low traffic (hops) to synchronize core and offloaded NDC, alias detection.

## Generality.

- Arbitrary combination of memory access pattern and compute operation.



# Existing NDC Abstractions

## Thread-Level NDC: Offload the entire thread context.




- Examples: TOM [ISCA '16], AMS [MICRO '18], ...
- ☒ Transparency.
- ☒ Synchronization-Efficiency.
- ☒ Generality: Works best with single data structure.

## Inst-Level NDC: Offload short instruction sequences.

- Examples: PIM-Enabled Inst [ISCA '15], Omni-Compute [ISCA '19], ...
- ☒ Transparency.
- ☒ Synchronization-Efficiency: At least one request per operation.
- ☒ Generality.


## User-Defined Function NDC: Offload user-defined operation.

- Examples: Active Routing [HPCA '19], Livia [ASPLOS '20], ...
- ☒ Transparency: Requires manually programming.
- ☒ Synchronization-Efficiency.
- ☒ Generality: Depends on the implementation.



# Streams: New Abstraction with Near-Data Comp.

```
while (i < N)  
  C[i] = A[i] + B[i]  
  i++
```



## **Stream: A decoupled seq. of addr/value.**


- Prevalent in data parallel workloads.
- Capture long-term patterns for each data structure  
→ Rich information for transparent and general offloading.

## **Near-Stream Computing:**

- Computation scheduled along with streams.

# Streams: New Abstraction with Near-Data Comp.

```
while (i < N)  
  C[i] = A[i] + B[i]  
  i++
```



## **Stream: A decoupled seq. of addr/value.**


- Prevalent in data parallel workloads.
- Capture long-term patterns for each data structure  
→ Rich information for transparent and general offloading.

## **Near-Stream Computing:**

- Computation scheduled along with streams.

# Streams: New Abstraction with Near-Data Comp.

```
while (i < N)
  C[i] = A[i] + B[i]
  i++
```



## Stream: A decoupled seq. of addr/value.

- Prevalent in data parallel workloads.
- Capture long-term patterns for each data structure  
→ Rich information for transparent and general offloading.

## Near-Stream Computing:

- Computation scheduled along with streams.
- ☒ Transparency: Extensive compiler/ISA support.
- ☒ Efficiency: Coarser-grained range-based sync.
- ☒ Generality: Support combinations of access patterns and ops.


## LLVM Compiler + GEM5 Simulator:

- **2.13×** speedup over SOTA NDC techniques.
- **76%** NoC traffic reduction.



# Outline

- Problems and Insights
- NDC Taxonomy and Opportunities
- Near-Stream Computing Implementation
- Evaluation



# Taxonomy of General Near-Data Computation

## Generality.

- Address pattern + compute type.

☐ *Affine*

☐ *Reduce*

☐ *Multi-Op*

☐ *Store*

☐ *Indirect*

☐ *RMW*

☐ *Ptr-Chase*

☐ *Load*

$A[i, j, \dots]$

$\sigma f(*A)$

$A[i], B[i], \dots$

$*A = f()$

$A[\dots B[i]]$

$*A = f(*A)$

$A = A.next$

$= f(*A)$

// Sum Array.

while (i < N)

  S += A[i]

  i++

// Vector Add.

while (i < N)

  C[i] = A[i] + B[i]

  i++

// Indirect RMW.

while (i < N)

  B[A[i]]++

  i++

// Link-List Op.

while (A)

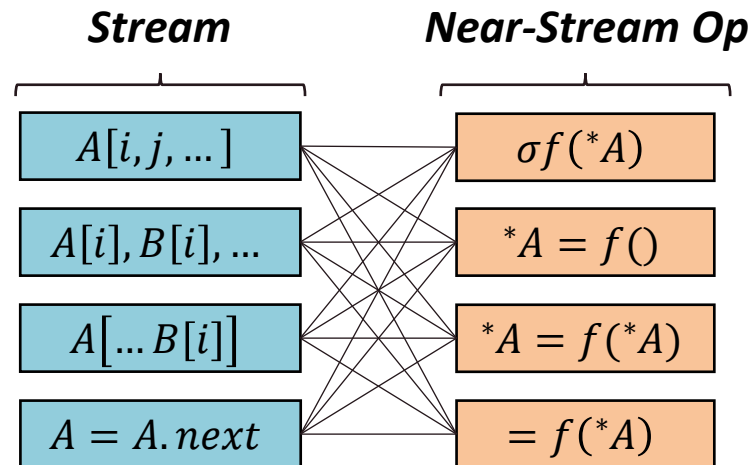
  f(A->value)

  A = A->next

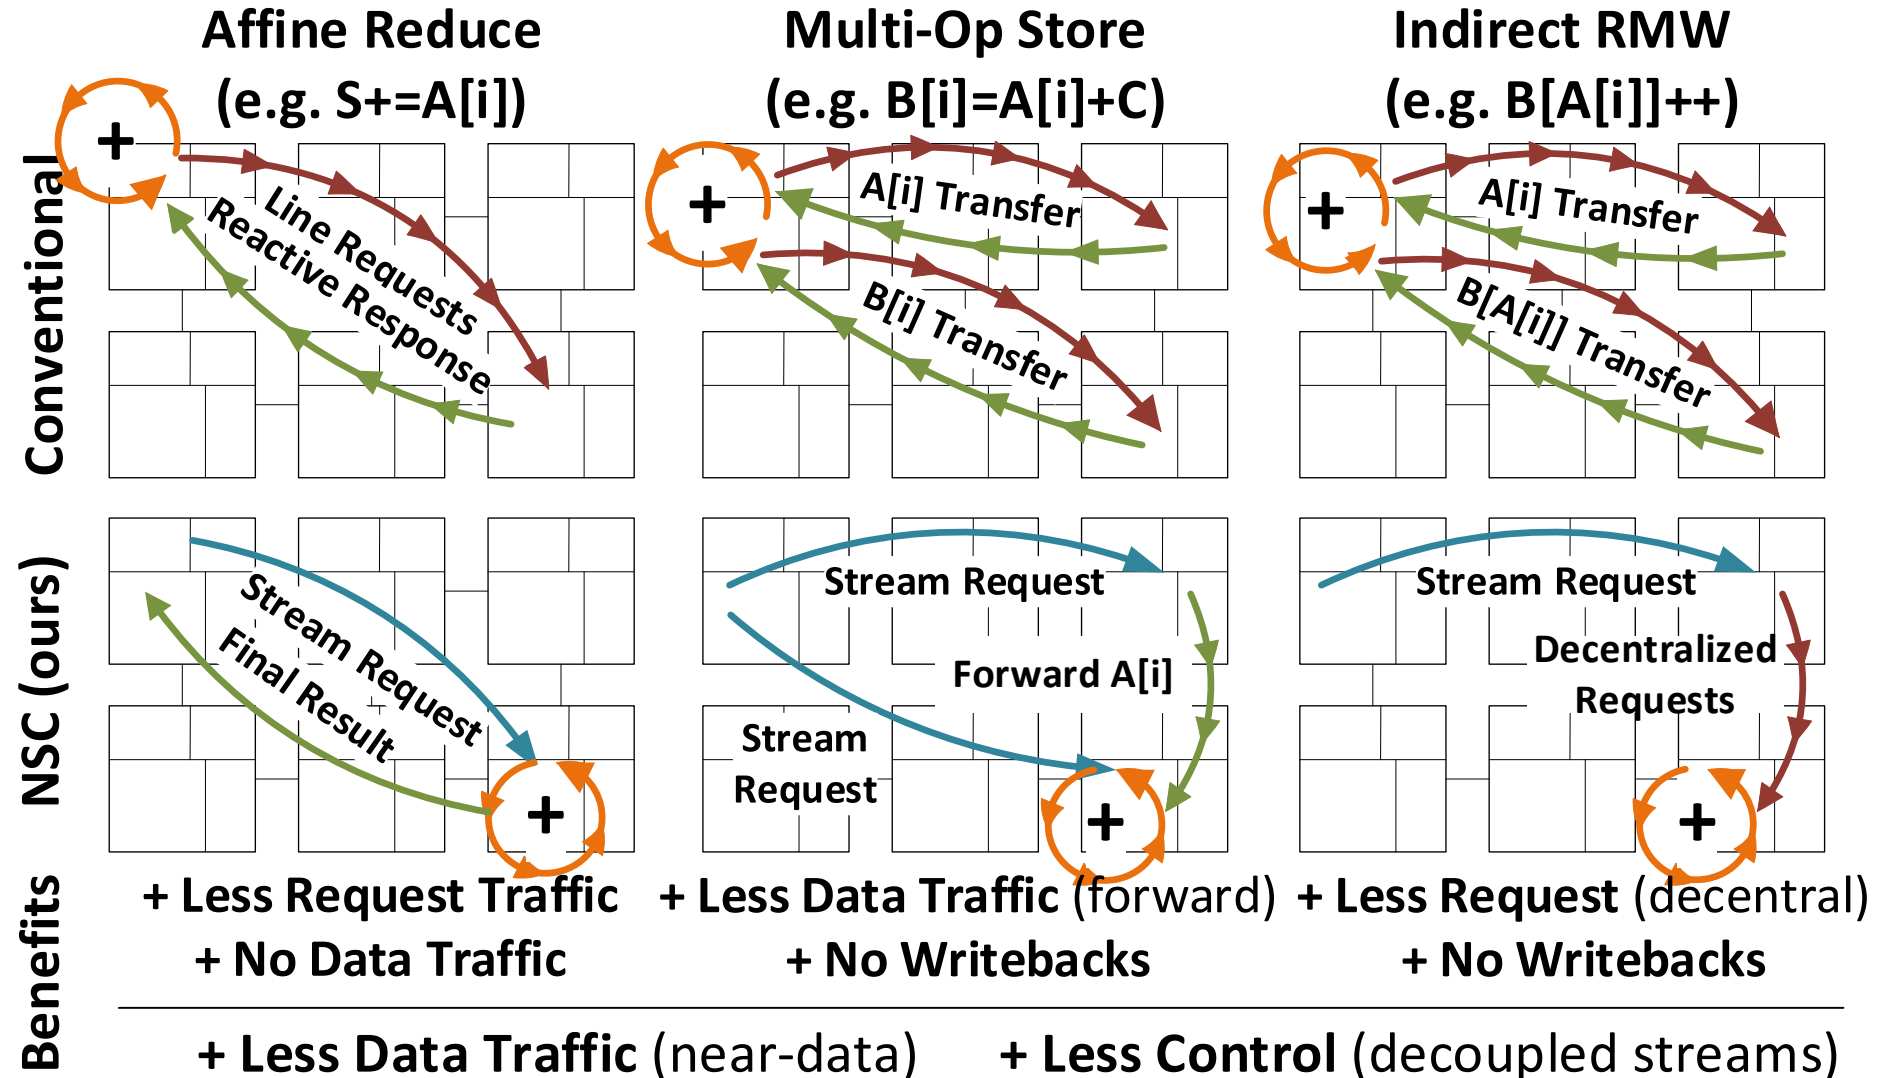
# Taxonomy of General Near-Data Computation

## Generality.

- Address pattern + compute type.
- ***Support all combinations.***
- Address Pattern → Streams; Compute Type → Near-Stream Operation.



# Conventional vs. Near-Stream Computing




# Breakdown of Near-Stream Computation

- 54% of dynamic micro-ops can be classified as near-stream computation.

**Major NS Comp.**


**Challenging to NDC:**

Inst-level NDC:  
always incurs  
high sync  
overheads.



# Outline

- Problems and Insights
- NDC Taxonomy and Opportunities
- Near-Stream Computing Implementation: How to ...
  - *Embed near-stream computing in the ISA?*
  - *Preserve sequential semantics?*
  - *Reduce coordination overheads?*
  - *Reuse hardware for computation?*
  - *When to offload computation?*
  - ...
- Evaluation



# Near-Stream Computing ISA Overview

- Explicitly encode streams and compute dependencies in the ISA.
- Compiler: Recognize streams and associated computation; Transform the program.
- **No manually programming required.**

*Original Pseudo Code*


*N-S Computing Pseudo Code*

*Stream Dep. Graph*

## (a) Vector Sum


```
while (i < N)
  v += A[i];
  i++;
```

```
s_cfg(sa=A[i], sv+=sa);
while (i < N)
  s_step(i);
v = s_load(sv);
s_end(sv, sa);
```



*The computation is outlined into a separate function, with the function pointer in the stream's configuration.*

# How to Compute with Low Overheads?




**Insight: There is already a core in the remote tile.**

Stream Compute Context (**SCC**):

- Light weighted thread context.
- No LSQ entries + Limited ROB entries.
- Looping over configured NSC functions.
- Released when streams are terminated.
- Reuse existing hardware units, e.g., decoder, vector units, ...



# Architectural Extensions



## **SE<sub>core</sub>:**

- Manage stream configuration and offload streams.
- Issue flow control credits and check for aliasing.

## **SE<sub>L3</sub>:**

- Fetch and forward operands to the consuming stream.
- Coordinated by SE<sub>core</sub> for flow control and aliasing check.


## **Stream Computing Manager (SCM):**

- Configure and terminate the stream computing context (**SCC**) with the near-stream function.
- Simple scalar ops handled by the ALU in SE<sub>core</sub> and SE<sub>L3</sub>  
Avoid frequently accessing SCC.

# Range-Sync: Coarse-Grained Sync. w. Seq. Semantics

**Challenges:** Efficiently synchronize the core and offloaded streams.

- Flow control scheme to avoid streams running too far ahead/behind.
- Detect possible core-stream and inter-stream aliasing.
- Provide sequential semantics for transparency and context switching.




# Range-Sync: Coarse-Grained Sync. w. Seq. Semantics

**Challenges:** Efficiently synchronize between core and streams.

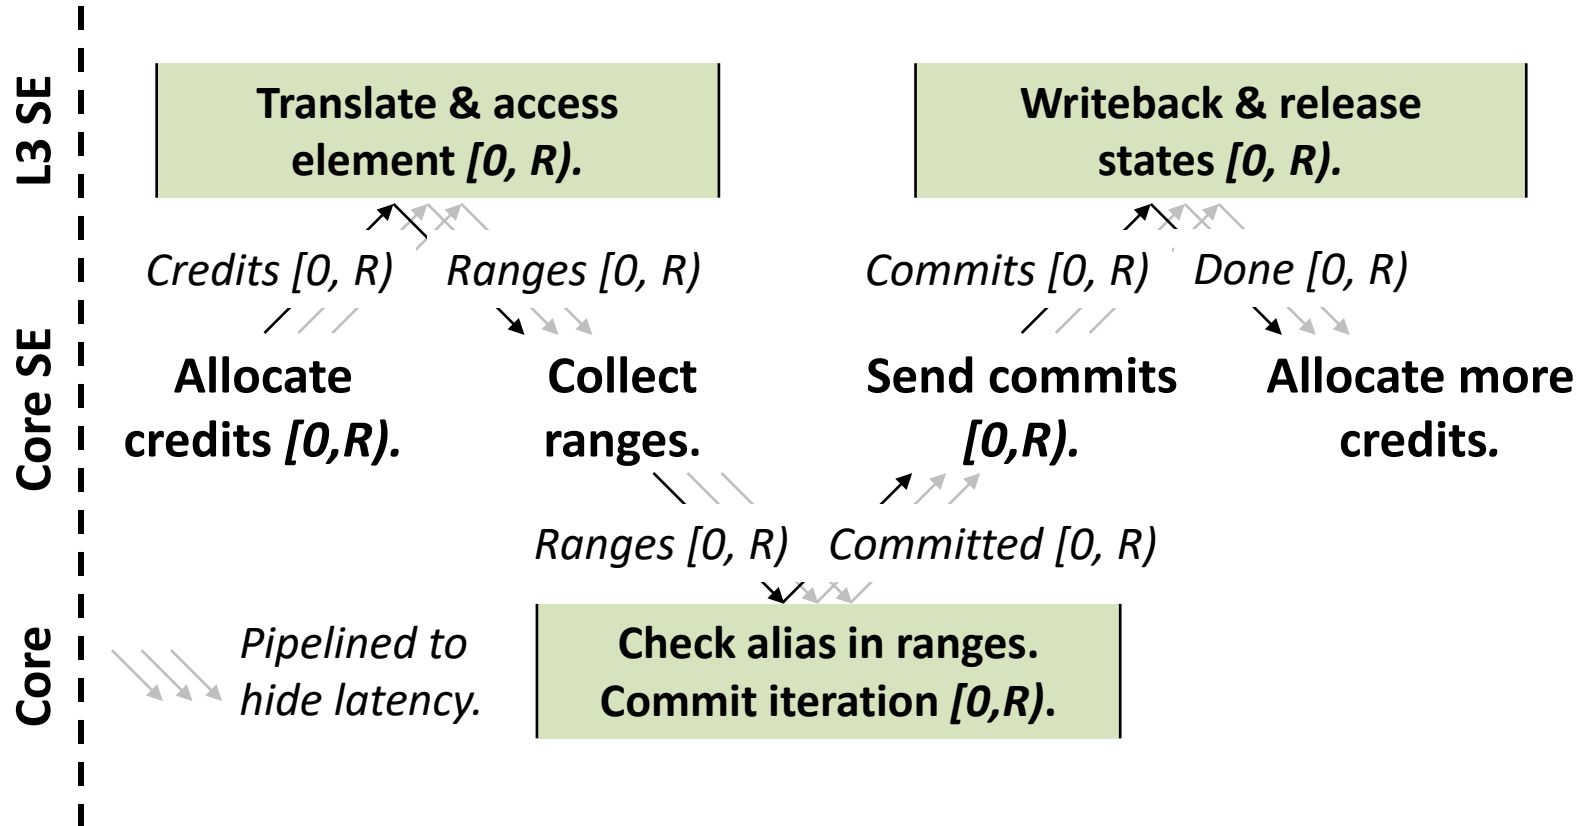
**Key Insight:** Sync. can be coarse-grained and conservative.

- Streams tend to access individual data structures, with non-overlapping addresses.
- Aliased streams should not be offloaded at all to avoid frequent synchronization.
- Allows false positives to dramatically reduce control overheads.

→ **Range-Sync:** Only sync. every few iterations against range of touched addresses.



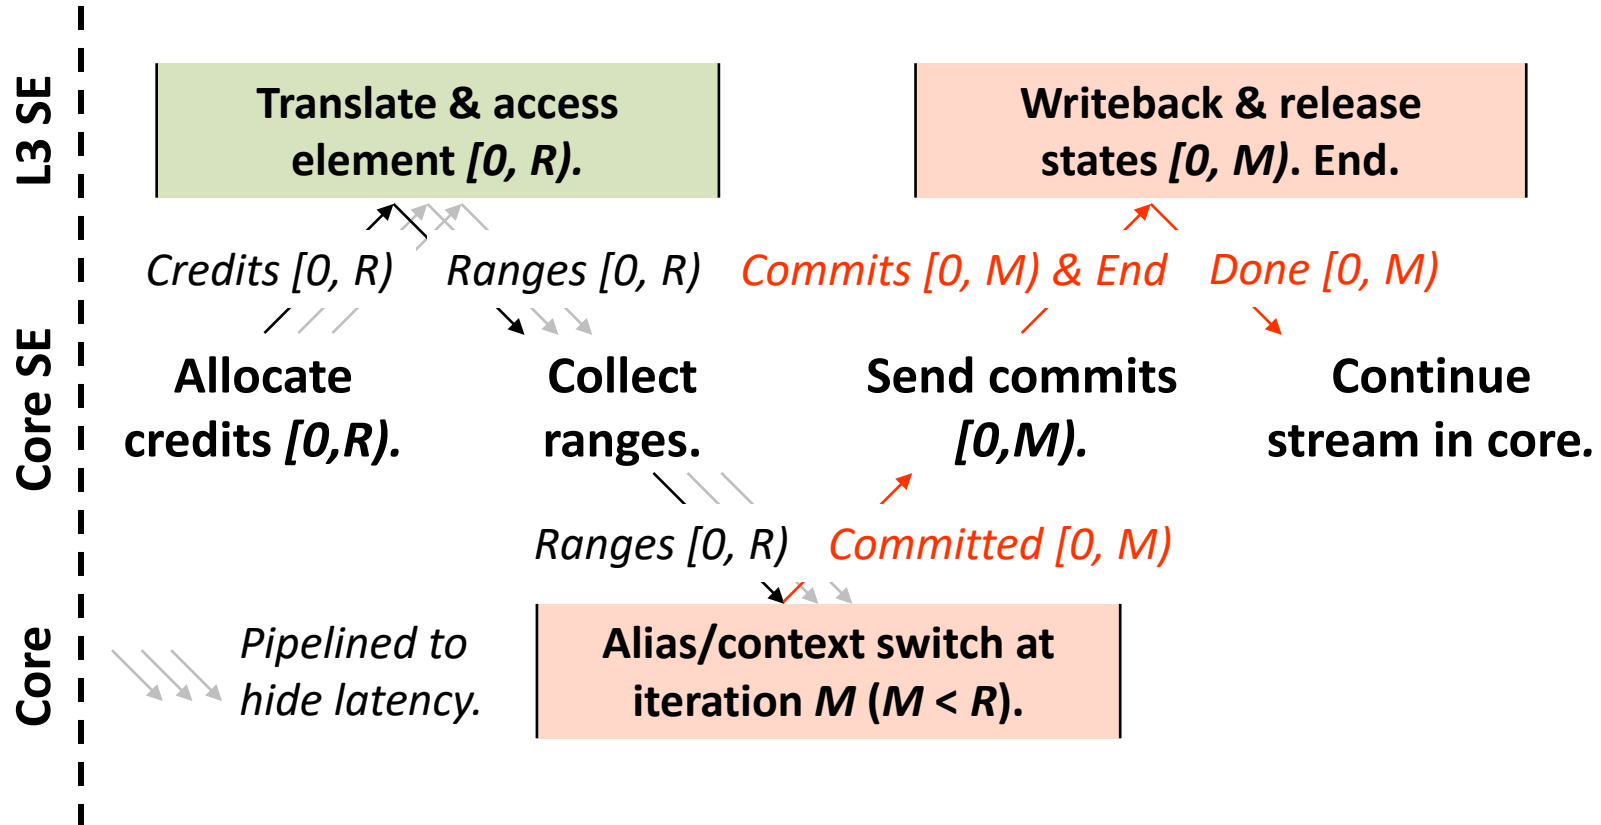
# Timeline of Range-Sync: Normal Case



## Key Takeaway:

- Sync every  $R$  iterations.
- Check alias against ranges.
- Overheads:  $4/R$  NoC msg. per affine element.

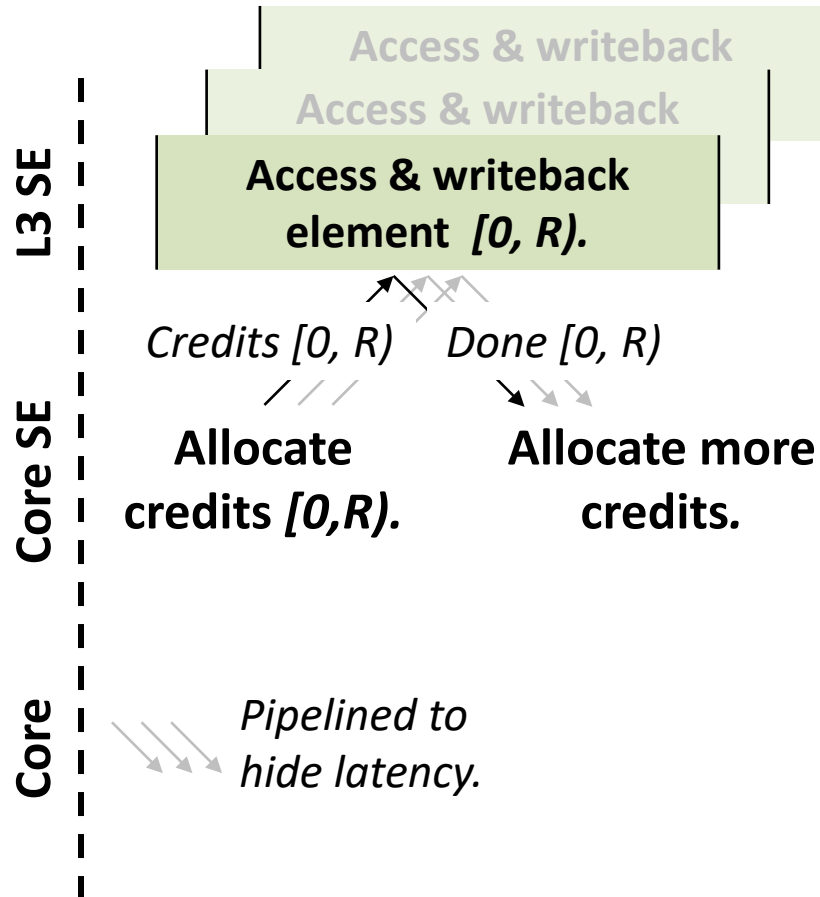
# Timeline of Range-Sync: Alias/Context Switch in Core



## Key Takeaway:

- *Execution continues at iteration  $M$  with streams in core.*
- *Still provide precise states.*

# Sync-Free Optimized NSC



## Key Takeaway:


- Pragma disables synchronization.
- Offloaded streams can directly commit.
- Coarse-grained context switch at iteration  $n \cdot R$ .
- Fully decoupled loops can be removed.
- → *Simultaneously execute multiple loops.*

## More details in the paper:

- Compiler transformation.
- Nested stream configuration.
- Avoid deadlocks.
- Determine when to offload.
- ...

# Outline

- Problems and Insights
- NDC Taxonomy and Opportunities
- Near-Stream Computing Implementation
- Evaluation



# Configurations


- LLVM-based compiler to recognize streams and transform programs.
- Gem5 20.0 cycle-level execution-driven simulator.
- 14 data processing workloads from Rodinia and Gap Graph Suite.
  - Parallelized with OpenMP, with AVX-512 enabled.
- Configurations (see paper for details):
  - 8x8 mesh topology, 3-level MESI, 32kB L1 I/D, 256kB L2, 1MB L3.
  - Base: Baseline cores with L1/L2 prefetchers. No stream support.
  - **Inst-Level NDC (INST)**: Computation is offloaded at instruction level to LLC.
  - **User-Def Func. NDC (UDF)**: Single operand UDF is offloaded at loop level.
  - $NS_{core}$ : Use  $SE_{core}$  as a prefetcher at the core [ISCA '19].
  - $NS_{no\ comp}$ : Streams are offloaded to LLC, but no computation [HPCA '21].
  - **NS**: Near-stream computations offloaded to LLC, with range-sync [this work].
  - **$NS_{no\ sync}/NS_{decouple}$** : Disable range-sync and fully decouple loops [this work].

} Resembles SOTA NDC:  
*Omni-Compute [ISCA '19]*  
*Livia [ASPLOS '20]*




# Overall Speedup with OOO8

- NS achieves 1.85x speedup over INST. (2.12x for NS<sub>decouple</sub> over UDF).
- Range-sync incurs low overheads as NS is only 8% slower than NS<sub>no sync</sub>.




# Dynamic Micro Ops Breakdown

- Generality: NSC captures 54% of dynamic micro-ops.
- 93% of possible ops are offloaded.




# NoC Traffic Breakdown

- INST and UDF reduces NoC traffic by 50% and 40% respectively.
- NS reduces NoC traffic by 69% (76% by NS<sub>decouple</sub>).
- NS offloads at stream granularity and incurs low overheads using range-sync.



# Energy vs. Speedup

- NS and NS<sub>decouple</sub> achieves 2.85×/3.52× energy efficiency over OOO8.
- Both small inorder cores and large OOO cores benefit from NSC.



# Conclusion

## Near-Stream Computing:

- ☒ Transparency: Extensive compiler/ISA support.
- ☒ Efficiency: Coarser-grained range-based sync.
- ☒ Generality: All combinations of address pattern + compute type.

**Core Idea: Encode access patterns and compute deps in ISA.**

→ Expose rich semantic information to the hardware.

→ Programmer-friendly near-data architectures.

